

Decisions Log & Known Issues / Lessons Learned

Consolidated file — merged from 011-decisions-log.md and 020-known-issues-lessons.md.

Source: 011-decisions-log.md (Decisions Log)

Decisions Log

Date	Decision	Rationale
2026-08-19	Scope = shift management only	User clarified mid-session; defer hospital modules
2026-08-19	hrms 16.5.0 pin	Lesson #44 — v16.5.1+ breaks on repost_allowed_types
2026-08-19	Shift code = 10-char [P][HHMM][S][HHMM]	User-proposed scheme, Option A (lean)
2026-08-19	Shift name = 10-char code itself (no separate shift_code field)	User simplification — name IS the code
2026-08-19	Holidays = standard Indian national + 4-5 Telangana	User confirmed
2026-08-19	Custom leave types = deferred	User said Haritha adds later
2026-08-19	Leave allocation = standard Indian defaults	User confirmed; rules in remarks column
2026-08-19	Source data = DO NOT modify	Canonicalization at import time only
2026-08-19	SSA is HRMS-native (Shift Schedule Assignment DocType)	Originally doubted; verified via docs.frappe.io/hr/shift-schedule-assignment
2026-08-19	Comprehensive 7-change schema update	9 HRMS docs verified; HRMS v15 canonical structure applied
2026-08-19	19 CSVs schema + data combined format	Manager-friendly for Google Sheets review
2026-08-19	3 designation collisions resolved automatically	Physician Asstant+Assistant, Sr.Executive+Senior Executive, Sr.Manager+Senior Manager
2026-08-19	3 shift code duplicates consolidated	A4+Shift-A, B2+Shift-B, C1+Shift-C
2026-08-20	Phase 0 + 1 signed off by manager	Schema + data CSVs approved
2026-08-20	New dedicated env pberp (clean slate)	Recommended over legacy envs
2026-08-20	Apps installed via bench install -app (runtime)	Quick start; lesson #47 trade-off (asset sync)
2026-08-20	Custom app: deferred, use custom fields + fixtures	Fast track for MVP

Date	Decision	Rationale
2026-08-20	Real-time employee name mapping: EMP-1001 (CSV) → HR-EMP-00001 (DB)	Built HR-EMP-{N-1000:05d} formatter
2026-08-20	Attendance imported via raw SQL (not ORM)	Bypassed Frappe Status validation + 240s timeout
2026-08-20	Attendance status options extended via Property Setter	Added “Weekly Off” + “Holiday” (1:1 CSV match)
2026-08-20	Employee Checkin via background jobs (25 batches × 500)	Direct console timed out on 12,562 rows
2026-08-20	HR-Attendance series counter fixed mid-flight (HR-ATT-2026- was 183, fixed to 6300)	Series was stale, prevented new record creation
2026-08-20	5 X-HH Department variants force-deleted via direct SQL	Frappe doc.delete() enforces “disable not delete”; direct DELETE bypasses
2026-08-21	Backend tested end-to-end via API: auth ☐ , CRUD ☐ , all 9 entities queryable ☐ , payroll/leave/holiday workflows ☐	
2026-08-21	UI smoke test inconclusive (headless browser tool unreliable)	Needs real browser verification next session
2026-08-21	Token limit issues (rate_limit_error) on long subagent runs — workaround: split into K1/K2/K3 + direct exec	Lesson learned for future orchestration
2026-08-21	nginx Upgrade: websocket forced for /socket.io/ (added during UI debugging — may need review/revert)	Frappe ws server validates Upgrade header on every request
2026-08-21	Rollback: pberp.duckdns.org env torn down 10:11-10:18 IST (Option B: nuke, no backup). All Phase 2-5 deployment work destroyed. Restart from Phase 1 on new env. CSV masters + git history intact.	Venkat authorized Option B at 10:33 IST; Phase 0 + 1 design work preserved; deployment was not recoverable
2026-08-25	Resumption: Phase 2 restart on pberpprod.duckdns.org (env created but never used = QA/dev effectively)	Wipe + reinit as Haritha. Same domain, clean slate. Backups-first mandatory (Step A1-A5).

Date	Decision	Rationale
2026-08-25	MEMORY correction: actual shift code format = [GMAN]\d{4}[RS]\d{4} (10-char)	Aug 19 decision said [P][HHMM][S][HHMM] (single P prefix, single S mid) — incomplete. Real data uses 4 prefixes (G/M/A/N) and R/S for end-time type. MEMORY.md needs update post-Phase 2.
2026-08-25	FK join key: shift_assignment.employee → em- ployee.attendance_device_id (EMP-NNNN)	Haritha employee.csv has NO name column (PK collision with HRMS). Use attendance_device_id as FK. 210/210 match verified.
2026-08-25	Subagent pattern rule: script work = subagent writes script + runs + fixes until works. Never hallucinate.	Applies to all scripted ops (backup, ingest, verify). Inline only for trivial ≤5-line edits.
2026-08-25	Subagent model selection: OX Alpha free for scripted ops	1M ctx, code-writing specialty, structured output. Nemotron 3 Ultra for reasoning-heavy. GLM 5.2 reserved (rate-limited).
2026-08-25	Verify script = reusable	scripts/verify_csvs.py re-runs in Phase 4 (CSV count vs DB count comparison).
2026-08-25	Phase 6 added: ISO/CMM Level 5 docs	Per Venkat 21:30 IST. Scope default = ISO 9001 + 27001, SOPs + process maps + audit trail, customer + manager audience.
2026-08-25	Phase 7 added: Handover + optional demo deck	Per Venkat 21:30 IST. Manager walkthrough + customer pilot (this week).
2026-08-25	Phase 2 Step A1-A5 backup executed + verified	Backup is mandatory before any destructive wipe (Aug 19 lesson #79 + SOUL never-migrate-prod-without- backup). OX Alpha subagent ran the script, parent (main) verified independently per Lesson #72.
2026-08-25	Git push: subagent fail-over to inline SSH	Nemotron 3 Ultra free returned FailoverError on first push attempt. Fell back to inline SSH chain (subagent quota/availability unreliable). Push succeeded: 219978d..7e95049 main -> main. 3 commits: f0e109a + 51d0bb4 + 7e95049. Branch now synced with origin.

Date	Decision	Rationale
2026-08-25	TRACKER.md updates via reusable script (update_tracker.py)	Per Venkat directive (22:46 IST): 'write a script for this as well because it is repetitive task'. JSON-driven, idempotent, handles status_date/footer/sections/decisions/lessons/p items. Future tracker updates = write JSON spec + run script. No more manual sed/edit.
2026-08-25	Subagent model fallback pattern established	Nemotron free FailoverError and OX Alpha rate limit both observed. Pattern: try primary model → on failure, fall back to inline or alternative free model. Document this for future ops. Premium MiniMax reserved for critical-path work.
2026-08-26	3d-1 Shift Assignment ingest needed Employees set to Active first	Frappe blocks 'Transactions cannot be created for an Inactive Employee'. All 210 Employees set to Active via per-row set_value (SQL UPDATE failed with column quoting bug).
2026-08-26	3d-2/3d-3 switched to raw SQL bulk insert (Lesson #43)	Attendance had 5 default status options but CSV uses 7 (incl. Holiday, Weekly Off). Property Setter fix didn't propagate to bench console session. ORM timed out on 12K+ checkin rows. Raw SQL bypasses both.
2026-08-26	Employee PK = HR-EMP-NNNNN (autoname), not CSV EMP-NNNN	HRMS Employee autoname uses naming series. Insert script must set employee_number, not name. CSV mapping: strip 'EMP-' prefix → employee_number → HR-EMP-NNNNN via DB query.
2026-08-26	Attendance status extended with Holiday + Weekly Off	CSV has 7 status values, Frappe default has 5. Property Setter added 'Holiday' and 'Weekly Off' to enable all CSV rows to insert.

Date	Decision	Rationale
2026-08-26	Synthetic-data defaults for required Employee fields	CSV has empty gender/date_of_birth (synthetic data). Defaults: gender='Not Specified' (created new Gender record), date_of_birth='1990-01-01', first_name=split(employee_name)[0].
2026-08-26	Phase 3.5 reconcile complete (Nemotron): all 11 entities match CSV targets	Department 47→37 (36 CSV + 'All Departments' root), Designation 76→48, Leave Type 9→7, Employment Type 6 (3 CSV-added Internship/Consultant/Temporary + 3 defaults), Holiday 28→14 (re-ingested). Lesson #72 parent-verify PASS on 11/11 entities.
2026-08-26	Bogus '(no rows)' Shift Location deleted	CSV ## Data section contained literal (no rows) placeholder. Ingest script did not check for this pattern, so it was inserted as a real record. Reconciler flagged + deleted. Pre-ingest must now detect this marker pattern.
2026-08-26	SS/SSA/SR synthesized (5 / 420 / 8)	SS = 5 templates (one per unique shift_type in SA rows), SR = 8 (status mix matched), SSA = 420 (one per unique employee × shift_type combo). All 5,318 SA rows linked via shift_assignment.shift_schedule_assignment FK. Script: scripts/synthesize_ssa_v2.py (commit 3f82928).
2026-08-26	HRMS SSA schema discovery: NO shift_type or date field	Brief schema said SSA has shift_type + date. Actual HRMS Shift Schedule Assignment is recurring template-bound — has only company, employee, shift_type, status, docstatus etc. NO date field. Original brief schema was wrong. SA rows link via shift_assignment.shift_schedule_assignment FK, not by date match.

Date	Decision	Rationale
2026-08-26	Cron regression at 10:06 IST Aug 26 — 3 backup lines dropped	Main VPS crontab now has only pberpprod_backup.sh line. dev_backup.sh, qa_backup.sh, erpclaw-git-daily-backup.sh were dropped (cause unknown — likely a crontab -e save gone wrong or an unrelated cron package reinstall). Streak 64/65 partial → 64/66 after slot 38 (18:00 IST Aug 26) missed. Restoring 3 lines gated on user YES/NO.
2026-08-27	Phase 3.6 bulk-submit needed 3 runs due to 3-layer Frappe framework barriers	Lesson #106: raw-SQL docs need naming_series backfill (#104), Property Setter adds meta.options but not controller-level status checks (#105), HRMS Attendance.validate() has hardcoded 5-value status list. Pattern applies to any submittable doctype bulk-submitted after raw-SQL ingest.
2026-08-27	Phase 3.6 scope was 6,314 docs, NOT 11,631 as estimated	Task brief assumed Shift Assignments were at docstatus=0. Reality: Phase 3.5 SSA synthesis script (commit 3f82928) submitted them as a side effect. User's list-view complaint was about Attendance + Holiday only. Always verify pre-state via Lesson #72 before estimating scope.
2026-08-27	Property Setter export → scripted recreate (Option 2)	Custom app fixtures not viable (PROD has no custom app on bench, apps/hrms/ is third-party core read-only per SOUL NEVER rule #3, manual JSON would need parallel applier). Script is idempotent, version-controlled, no core edits, reusable on any env. Mirrors bulk_submit.py's importlib.util.spec_from_file_location() pattern.

Source: 020-known-issues-lessons.md (Known Issues / Lessons Learned)

Known Issues / Lessons Learned

#	Issue	Lesson
47	Asset sync after bench build requires per-directory docker cp with /. syntax	Backend + frontend have separate /sites/assets volumes
80	Frappe installed_apps lives in TWO places: site_config.json AND sites/apps.txt	Both must be edited atomically
44	HRMS v16.5.1+ install broken by repost_allowed_types phantom field	Pin to v16.5.0
79	bench backup --with-files has no built-in timeout	Always wrap with timeout 900 + capture \${PIPESTATUS[0]}
NEW	bench install-app runtime creates asset hash drift (Lesson #47)	For future envs, bake apps into image via apps.json
NEW	WebSocket container can crash with Redis SocketClosedUnexpectedlyError	Add health check + auto-restart
NEW	Frappe ws server requires Upgrade: websocket header even for polling	nginx needs explicit proxy_set_header Upgrade "websocket"; for /socket.io/
NEW	Frappe doc.delete() enforces "disable not delete"	Force-delete via direct DB DELETE
NEW	Token limit errors on long subagent runs (rate_limit_error)	Split large tasks into smaller subagents OR direct exec
NEW	Frappe headless browser testing unreliable (timeouts, false negatives)	Use real browser for UI smoke tests
NEW	Subagent work pattern: write script + run + fix loop until works; never hallucinate (2026-08-25)	Applies to all scripted ops. Inline only for trivial ≤5-line edits.
NEW	Verify scripts must discover ## Data marker correctly (2026-08-25)	csv.DictReader needs header in input — slice readlines()[header_idx:] (marker+1), not [data_start:] (marker+2).
NEW	Haritha employee.csv has no name column (2026-08-25)	Use attendance_device_id (EMP-NNNN) as FK join key for shift_assignment.employee.

#	Issue	Lesson
OX Alpha free model rate-limited mid-task (Aug 25 commit subagent)	Subagent bootstrap costs ~13-15k tokens even when LLM call is free. For trivial mechanical ops (git commit, single SSH call), inline is cheaper + faster. Reserve subagent for scripted ops with fix loops (write script → run → fix until works).	
Nemotron 3 Ultra free returned FailoverError mid-push (Aug 25)	Free-tier models have unreliable availability. For critical ops (e.g., git push to protected branch), have inline fallback ready. Don't depend on subagent success for one-shot operations.	
Subagent claimed success on backup, parent verification needed (Lesson #72)	Always run independent verification probes after subagent claims (SHA256 byte-match, file listing, gzip integrity). Costs ~500 tokens inline; saves catching fabricated success reports.	
scp directly from gateway to venkat VPS failed (Aug 25 backup)	Backup files on vijay VPS, not gateway. Use <code>ssh tar-pipe</code> from vijay to venkat for cross-VPS copy, or move files through a shared mount. Don't assume direct paths between VPS hosts.	
TRACKER.md manual edits are repetitive + error-prone	Use <code>scripts/update_tracker.py</code> (Aug 25) — JSON-driven, idempotent. Future updates = JSON spec + run script. Covers: <code>status_date</code> , footer, sections, decisions log, lessons, pending actions.	
Frappe 16 Property Setter changes don't refresh bench console meta cache	After setting Property Setter for Select options, need to open NEW bench console session (or restart workers) for new options to take effect. Within same session, meta is cached and old options list is used even after <code>frappe.clear_cache()</code> .	

#	Issue	Lesson
MariaDB UPDATE with backticks around table name + plain column name failed: 'Unknown column "Active" in SET'	Use <code>frappe.db.set_value()</code> per-row for safety, or check exact column quoting. SQL string escaping is finicky across Frappe/Python/MariaDB combinations.	
ipython cell splitting breaks multi-statement scripts via <code>exec(open())</code>	ipython/bench console interprets heredoc input as multiple cells (split at blank lines / function defs). Variables defined in one cell aren't accessible in another. Workaround: use single-line code or wrap everything in <code>main()</code> function called from one cell.	
Frappe HRMS Shift Schedule Assignment is recurring template-bound — has NO <code>shift_type</code> or <code>date</code> field	Link SA rows via the <code>shift_assignment.shift_schedule_assignment</code> FK field, not via <code>date/shift_type</code> matching. Original SSA draft tried to create one SSA per employee × <code>shift_type</code> × <code>date</code> combo (over-counted to 1,758); correct is one SSA per unique employee × <code>shift_type</code> (420). Verify HRMS DocType JSON before drafting brief.	
Auto-name with <code>autoname = field:source</code> does NOT enforce uniqueness on display name	Multiple rows can share the same <code>department_name</code> (or any other display field) while having unique name PKs. Reconciling requires grouping by display name + keeping oldest (or matching CSV target count). Lesson #73 pattern: always diff <code>COUNT(*)</code> vs CSV row count before declaring master ingest done.	
(no rows) literal placeholder in CSV ## Data section will be ingested as a real record	CSV empty-marker convention (some tools emit (no rows) instead of zero data lines) must be detected by pre-ingest script. Insert one naive line and you get a bogus master record (e.g., Shift Location named '(no rows)'). Reconciler caught this on Phase 3.5 — add explicit check in <code>scripts/ingest_masters.py</code> for any (no rows) or <empty> literal pattern.	

#	Issue	Lesson
104	Raw SQL ingest bypasses Frappe's mandatory field defaults — submit() fails on reqd=1 fields that auto-set during ORM insert (Phase 3.6, 2026-08-27)	Docs inserted via raw SQL (bypassing Frappe ORM) don't populate fields set by validate() or before_insert() hooks (e.g., naming_series on Attendance). When submit() runs, it re-validates and fails on the missing mandatory field. Workaround: UPDATE tabX SET <field>=<series> WHERE docstatus=0 AND (<field> IS NULL OR <field>='') before bulk-submit.
105	Property Setter for Select options does NOT bypass controller-level hardcoded status checks (Phase 3.6, 2026-08-27)	Adding 'Holiday' and 'Weekly Off' to Attendance's status options via Property Setter is necessary but NOT sufficient. HRMS Attendance.validate() calls erp-next.controllers.status_updater.validate_status(...) with a hardcoded 5-value list at apps/hrms/hrms/hr/dctype/attendance/attendance.py. The controller-level check rejects values not in the hardcoded list, even if Property Setter has added them. Fix: monkey-patch erp-next.controllers.status_updater.validate_status before submit() — wrap to silently accept the extra statuses. Cannot fix at Property Setter level without editing HRMS core (SOUL NEVER rule).
106	Bulk-submit of raw-SQL-inserted submittable docs may need 3 runs: (1) backfill mandatory fields, (2) add Property Setter, (3) monkey-patch controller-level checks (Phase 3.6, 2026-08-27)	Pattern observed for Haritha Attendance: 6,300 docs needed all three fixes. Each fix is fast (~10-60s for the run) but cumulatively adds 2 extra runs. Plan for 3 runs in time estimates. Total wall time: ~10 min for 6,300 docs on pberpprod.

#	Issue	Lesson
frappe.make_property_setter has TWO implementations with DIFFERENT signatures (Phase 3.7, 2026-08-27)	Property Setter frappe.make_property_setter(args_dict, ignore_validate=False, validate_fields_for_doctype=True, is_system_generated=True, *, module=None) takes a dict-like args. Lower-level frappe.custom.doctype.property_setter.property_setter.make_property_setter(doctype, fieldname, property, value, property_type, for_doctype=False, validate_fields_for_doctype=True, is_system_generated=True) uses positional + for_doctype kwarg. The dict version does NOT accept for_doctype — it derives doctype_or_field from args.doctype_or_field (default 'DocField'). Calling the wrong signature raises TypeError: unexpected keyword argument.	
bench export-fixtures silently skips Property Setters for apps whose hooks.py doesn't list them (Phase 3.7, 2026-08-27)	export-fixtures iterates frappe.get_hooks('fixtures', app_name=app) per app. If hooks.py has no fixtures = [...] list containing 'Property Setter', the export produces no output (no error, no warning, exit 0). Property Setters live in DB only and don't migrate on env rebuild / bench update. For Haritha's HRMS Property Setter: hooks.py is third-party code (SOUL NEVER rule #3 forbids editing). Solution: scripted recreate, not fixture export.	
bench execute requires module to be importable from cwd — /tmp/ scripts fail (Phase 3.7, 2026-08-27)	frappe.get_attr() first segment must be an installed app name (raises AppNotInstalledError otherwise). Fallback eval(code) needs the module already imported. Use bench console < /tmp/wrapper.py + importlib.util.spec_from_file_location() pattern (same as bulk_submit.py). The wrapper imports the /tmp/ script as a module then calls its run() function.	

#	Issue	Lesson
112	Raw SQL bulk INSERT bypasses ORM auto-derivation of FK fields (Phase 3.9, 2026-08-27)	Frappe ORM <code>frappe.get_doc().insert()</code> auto-derives FK-derived fields like <code>employee_name</code> and <code>department</code> from the linked parent DocType at insert time. Raw SQL <code>INSERT INTO tabAttendance (...) VALUES (...)</code> does NOT — the FK-derived columns end up NULL/empty. Same root cause as Lesson #104 (<code>naming_series</code>) and Lesson #110 (linkage fix). Future raw-SQL ingest scripts must either enumerate derived fields explicitly in the column list, OR plan a post-ingest populate script. Lesson #111's pattern of "verify pre-state then post-state counts" applies here too.
<code>frappe.db.sql()</code> returns empty tuple for UPDATE statements in MariaDB (Phase 3.9, 2026-08-27)	<code>cursor.rowcount</code> carries the affected-row count, but the SQL result tuple is <code>()</code> . Code that does <code>result = frappe.db.sql("UPDATE ..."); matched = result[0][0]</code> always reads 0. Workaround: compute the diff between pre-state and post-state counts, OR check <code>frappe.db._cursor.rowcount</code> directly. Affects idempotency verification — re-running UPDATE on already-populated rows will return <code>matched=0</code> from SQL but show actual 0 rows changed only if you verify via WHERE-matched count separately.	
<code>frappe.db.commit()</code> after raw SQL UPDATE is required (Phase 3.9, 2026-08-27)	MariaDB connector default is <code>autocommit OFF</code> in bench. UPDATE rows are only persisted after explicit <code>frappe.db.commit()</code> . Without it, a follow-up <code>frappe.db.sql("SELECT COUNT(*)")</code> from a fresh bench console session will still see the old state. Same pattern as Phase 3.8 (Lesson #111).	

Last updated: Phase 3.7 idempotent `recreate_property_setters.py` added (15:06 IST). Fixes Rule #9 violation — Attendance-status-options Property Setter was DB-only, no fixture export possible (HRMS `hooks.py` doesn't list Property Setter). Tested: delete PS → run script → PS recreated → run again → no duplicate. Script path: `scripts/recreate_property_setters.py`. Cron

regression also resolved (commit 5f383b6, 4 backup lines now active: dev/qa/pberpprod/git-daily).
